

```
In [14]: # sanity checks
import sys, platform, pkgutil
import pandas as pd
print("Python:", sys.executable)
print("Version:", platform.python_version())
print("xarray installed?", pkgutil.find_loader("xarray") is not None)
```

Python: /Users/vb/Desktop/VBThesisCode2_FastAPI/.venv/bin/python
Version: 3.11.11
xarray installed? True

```
In [15]: import xarray as xr
# path for the NetCDF file downloaded from Davis weather station data at
nc_path = "../input/ExternalTemp/davis-TUD-GV_Green_Village_202503.nc"

ds = xr.open_dataset(nc_path)
ds # print a summary: dimensions, coords, data_vars, attributes
```

Out [15]: xarray.Dataset

► Dimensions: (time: 44636)

► Coordinates: (0)

► Data variables:
(16)

► Indexes: (0)

► Attributes: (11)

```
In [16]: # List variables & attributes - to find temperature name/units
print("Variables:", list(ds.data_vars))
for v in ds.data_vars:
    print(f"\n[{v}]")
    for k, val in ds[v].attrs.items():
        print(f"  {k}: {val}")
```

Variables: ['epoch_time', 'time_as_string', 'latitude', 'longitude', 'altitude', 'pressure', 'temperature', 'humidity', 'solar_irradiance', 'rain', 'rain_rate', 'UV_radiation', 'wind_speed', 'wind_to_direction', 'wind_gust_speed', 'wind_gust_from_direction']

[epoch_time]

long_name: End time of observation period, in seconds since 1970-01-01 00:00:00 UTC

units: s

comment: The length of a measurement interval is 1 minute. Values are rounded to the nearest integer.

[time_as_string]

long_name: End time of observation period in YYYY-MM-DD HH:MM:SS UTC

units:

comment: This is the same as epoch_time but in human readable format.

[latitude]

long_name: latitude of site

units: degree_north

[longitude]

long_name: longitude of site

units: degree_east

[altitude]

long_name: altitude of site

units: m

[pressure]

long_name: atmospheric pressure

units: mbar

comment: inside the console

[temperature]

long_name: air temperature

units: °C

comment: outside, within the radiation shield

[humidity]

long_name: relative humidity

units: %

comment: within the radiation shield, rounded to 1%

[solar_irradiance]

long_name:

units: W/m2

[rain]

long_name: rainfall

units: mm

comment: determined using a 0.2 mm tipping bucket

[rain_rate]

long_name: rainfall rate

units: mm/h

comment: calculated using the WeeWX RainRater parameter.

[UV_radiation]

long_name: UV radiation

units: W/m2

[wind_speed]

long_name: wind speed

units: m/s

comment: averaged over the measurement period

[wind_to_direction]

long_name: wind to direction

units: degrees

comment: clockwise from North

[wind_gust_speed]

long_name: wind gust speed

units: m/s

[wind_gust_from_direction]

long_name: wind gust from direction

units: degrees

comment: clockwise from North

```
In [17]: # Convert to a tidy DataFrame and preview
df = ds.to_dataframe().reset_index()
df.head()
# df.columns
```

```
Out[17]:
```

	time	epoch_time	time_as_string	latitude	longitude	altitude	pressure
0	0	1740787200	2025-03-01	51.996109	4.378741	2.0	1027.328613
1	1	1740787260	2025-03-01 00:01:00	51.996109	4.378741	2.0	1027.328613
2	2	1740787320	2025-03-01 00:02:00	51.996109	4.378741	2.0	1027.328613
3	3	1740787380	2025-03-01 00:03:00	51.996109	4.378741	2.0	1027.328613
4	4	1740787440	2025-03-01 00:04:00	51.996109	4.378741	2.0	1027.328613

```
In [18]: # non-NaN temperature values
non_nan_count = df['temperature'].notna().sum()
print(f"Non-NaN temperature values: {non_nan_count}")

# show total rows and percentage
total_rows = len(df)
percentage = (non_nan_count / total_rows) * 100
print(f"Total rows: {total_rows}, Percentage non-NaN: {percentage:.2f}%")
```

Non-NaN temperature values: 44027

Total rows: 44636, Percentage non-NaN: 98.64%

Creating a new dataframe that has columns -

1. pressure,
2. temperature,
3. humidity,

4. solar_irradiance,
5. wind_speed.

For each of these columns, storing only one value per hour.

For Jan 2025 data - the 'time_as_string' column has 31 NaN values in 43201. Not a big problem but epoch time has been used instead (since that has no NaN values) epoch_time values are Unix seconds (increments of 60 = 1-minute steps), so we should convert with unit='s'

```
In [21]: # CONVERTING epoch_time to datetime format
# REASON - df.resample('1h') only works if the DataFrame's index is a Dat

# --- 1) Build a UTC timestamp from epoch_time (seconds since 1970) ---
# If epoch_time ever comes as strings, this coerces to numeric first.
df["timestamp_utc"] = pd.to_datetime(pd.to_numeric(df["epoch_time"], erro

bad = df["timestamp_utc"].isna().sum()
if bad:
    print(f"Dropping {bad} rows with bad epoch_time")
clean_df = df.dropna(subset=["timestamp_utc"])

# --- 2) Index by time and sort ---
clean_df = clean_df.set_index("timestamp_utc").sort_index()

# --- 3) Choose sensor columns ---
cols = ["pressure", "temperature", "humidity", "solar_irradiance", "wind_

# Coerce selected columns to numeric (strings -> numbers; invalid -> NaN)
clean_df[cols] = clean_df[cols].apply(pd.to_numeric, errors="coerce")

# --- 4) Resample to hourly means (NaNs ignored; hour with all NaNs remai
clean_df_hourly = clean_df[cols].resample("1h").mean()

# Back to a regular column + add local time if useful
clean_df_hourly = clean_df_hourly.reset_index()
clean_df_hourly["time_local"] = clean_df_hourly["timestamp_utc"].dt.tz_co

# Optional: keep only hours with a valid temperature
# df_hourly = df_hourly.dropna(subset=["temperature"])

clean_df_hourly.head()
```

Out [21]:

	timestamp_utc	pressure	temperature	humidity	solar_irradiance	wind_sp
0	2025-03-01 00:00:00+00:00	1027.328613	2.902500	89.016667	0.0	0.000
1	2025-03-01 01:00:00+00:00	1027.328613	3.022833	89.850000	0.0	0.000
2	2025-03-01 02:00:00+00:00	1027.328613	3.227333	90.000000	0.0	0.000
3	2025-03-01 03:00:00+00:00	1027.328613	3.878667	90.000000	0.0	0.000
4	2025-03-01 04:00:00+00:00	1027.328613	3.920167	89.533333	0.0	0.000

Note - For Jan 2025, most of the visible rows had NaN as the hourly average for sensor data (temperature, humidity, solar_irradiance and wind_speed). This is not the case in Feb 2025. But still, looking further into the extent NaN values - first wrt all sensor data and then specifically wrt temperature data

```
In [22]: # List the sensor columns
sensor_cols = ["temperature", "humidity", "solar_irradiance", "wind_speed"]

# Count of rows where at least ONE sensor value is NaN
at_least_one_nan_count = clean_df_hourly[sensor_cols].isna().any(axis=1).sum()
print(f"Number of hours where at least ONE sensor value is NaN: {at_least_one_nan_count}")

# Find rows where ALL sensor values are NaN
nan_mask = clean_df_hourly[sensor_cols].isna().all(axis=1)

# Count them
nan_count = nan_mask.sum()
print(f"Number of hours with all NaN sensor values: {nan_count}")

# Create a new dataframe with only those rows
nan_rows_df = clean_df_hourly[nan_mask]

# Count of unique days in nan_rows_df
unique_days_with_all_nan = nan_rows_df["time_local"].dt.date.nunique()
print(f"Number of unique days with all NaN hours: {unique_days_with_all_nan}")

# Show the first few for inspection
print(nan_rows_df.head())

# (Optional) check percentage
total_rows = len(clean_df_hourly)
print(f"Percentage: {100 * nan_count / total_rows:.2f}%")
```

Number of hours where at least ONE sensor value is NaN: 1

Number of hours with all NaN sensor values: 0

Number of unique days with all NaN hours: 0

Empty DataFrame

Columns: [timestamp_utc, pressure, temperature, humidity, solar_irradiance, wind_speed, time_local]

Index: []

Percentage: 0.00%

Temporal distribution of all NaN rows

```
In [23]: import matplotlib.pyplot as plt
import matplotlib.dates as mdates

nan_mask = clean_df_hourly[sensor_cols].isna().all(axis=1)

fig, ax = plt.subplots(figsize=(14, 4)) # taller & wider


# Plot all-NaN hours as vertical lines
ax.vlines(
    clean_df_hourly.loc[nan_mask, "time_local"],
    ymin=0, ymax=1,
    color="red", alpha=0.8, linewidth=2, label="All sensors NaN"
)

# Plot hours with some data present
ax.vlines(
    clean_df_hourly.loc[~nan_mask, "time_local"],
    ymin=0, ymax=0.2,
    color="green", alpha=0.4, linewidth=1, label="Some data present"
)

# Formatting
ax.set_ylim(-0.1, 1.1)
ax.set_yticks([]) # hide y-axis ticks
ax.set_title("Hours with All NaN Sensor Values (Vertical Lines)")
ax.set_xlabel("Time")

# Better date formatting: days as major ticks, hours as minor
ax.xaxis.set_major_locator(mdates.DayLocator(interval=1))
ax.xaxis.set_major_formatter(mdates.DateFormatter("%b %d"))
ax.xaxis.set_minor_locator(mdates.HourLocator(byhour=[0, 6, 12, 18]))
ax.xaxis.set_minor_formatter(mdates.DateFormatter("%H:%M"))

plt.setp(ax.xaxis.get_majorticklabels(), rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```



When do gaps occur? — histograms

By hour-of-day (helps spot diurnal patterns): NOT FOR MARCH 2025 AS NO ROWS WITH ALL NAN DATA

```
In [ ]: nan_by_hour = nan_rows_df["time_local"].dt.hour.value_counts().sort_index
ax = nan_by_hour.plot(kind="bar", figsize=(10,3))
ax.set_title("All-NaN Hours by Hour-of-Day")
ax.set_xlabel("Hour of day")
ax.set_ylabel("Count of all-NaN hours")
plt.tight_layout()
plt.show()
```

By day (helps see clustering by dates): NOT FOR MARCH 2025 AS NO ROWS WITH ALL NAN DATA

```
In [ ]: nan_by_day = nan_rows_df["time_local"].dt.date.value_counts().sort_index()
ax = nan_by_day.plot(kind="bar", figsize=(12,3))
ax.set_title("All-NaN Hours per Day")
ax.set_xlabel("Date")
ax.set_ylabel("Count of all-NaN hours")
plt.tight_layout()
plt.show()
```

Focusing on NaN values under 'Temperature' column

```
In [26]: # ==== Temperature gap analysis using clean_df_hourly (expects time_local)

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.dates as mdates

dfh = clean_df_hourly[['time_local', 'temperature']].dropna(subset=['time
dfh = dfh.sort_values('time_local').reset_index(drop=True)

# Mask of missing temperature
temp_na = dfh['temperature'].isna()

# ----- Summary
total_hours = len(dfh)
missing_hours = int(temp_na.sum())
pct_missing = 100.0 * missing_hours / total_hours if total_hours else 0.0
print(f"Total hours: {total_hours} | Missing temp hours: {missing_hours}")

# ----- Continuous missing segments (start, end, duration)
grp_id = (temp_na != temp_na.shift()).cumsum()
runs = (
    dfh.assign(_na=temp_na, _grp=grp_id)
    .groupby('_grp', as_index=False)
    .agg(is_na=('_na', 'first'),
         start=('time_local', 'first'),
         end=('time_local', 'last'),
         n_hours=('time_local', 'size'))
)
gap_runs = runs[runs['is_na']].drop(columns='is_na').reset_index(drop=True)
```

```

if not gap_runs.empty:
    longest = gap_runs.iloc[gap_runs['n_hours'].idxmax()]
    print("Longest gap:", longest['start'], "→", longest['end'], f"({int(
else:
    print("No continuous missing segments.")

display(gap_runs.head(10))

# ----- Visualization A: timeline of missing hours
fig, ax = plt.subplots(figsize=(12, 3.5))
ax.vlines(dfh.loc[temp_na, 'time_local'], ymin=0, ymax=1, linewidth=2, al
ax.set_ylim(-0.05, 1.05); ax.set_yticks([]); ax.set_title("Temperature -
ax.set_xlabel("Time")
ax.xaxis.set_major_locator(mdates.DayLocator(interval=1))
ax.xaxis.set_major_formatter(mdates.DateFormatter("%b %d"))
ax.xaxis.set_minor_locator(mdates.HourLocator(byhour=[0,6,12,18]))
plt.setp(ax.xaxis.get_majorticklabels(), rotation=45)
plt.tight_layout(); plt.show()

# ----- Visualization B: day×hour heatmap (1=missing, 0=present)
grid = (pd.DataFrame({
    'date': dfh['time_local'].dt.date,
    'hour': dfh['time_local'].dt.hour,
    'is_missing': temp_na.astype(int)
}))
    .groupby(['date', 'hour'], as_index=False)['is_missing'].max())

pivot = grid.pivot(index='date', columns='hour', values='is_missing').sor

fig, ax = plt.subplots(figsize=(12, max(3, 0.25*len(pivot))))
im = ax.imshow(pivot.values, aspect='auto', interpolation='nearest')
ax.set_title("Temperature Missingness (1=missing, 0=present)")
ax.set_xlabel("Hour of Day"); ax.set_ylabel("Date")
ax.set_xticks(np.arange(0,24,3)); ax.set_xticklabels([str(h) for h in ran
ax.set_yticks(np.arange(len(pivot.index)))
ax.set_yticklabels(pd.to_datetime(pivot.index).strftime("%b %d"), fontsiz
cbar = plt.colorbar(im, ax=ax, fraction=0.046, pad=0.04)
cbar.set_ticks([0,1]); cbar.set_ticklabels(['present', 'missing'])
plt.tight_layout(); plt.show()

```

Total hours: 745 | Missing temp hours: 1 (0.13%)


Longest gap: 2025-03-04 07:00:00+01:00 → 2025-03-04 07:00:00+01:00 (1 h)

	_grp	start	end	n_hours
0	2	2025-03-04 07:00:00+01:00	2025-03-04 07:00:00+01:00	1

[illegible]

Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 6425 ticks ([19348.25, ..., 20954.25]), which exceeds Locator.MAXTICKS (1000).
Locator attempting to generate 1606 ticks ([19349.0, ..., 20954.0]), which exceeds Locator.MAXTICKS (1000).

Temperature — Hours with Missing Values



Time

